

Report on the estimation of normalization coefficient for diffusion based method with deep convolutional neural networks

Guillaume Coulaud

August 28, 2022

Contents

1	Introduction	3
2	Problem definition	4
3	Choosing the architecture	6
4	Finding a good baseline	7
4.1	Distance map from coasts	7
4.2	Activation function	8
5	Skip connection	10
5.1	First dataset	10
5.2	Change of dataset	11
5.3	First architecture	11
5.4	Second architecture	13
6	Data augmentation	16
6.1	Method	16
6.2	Results	17
7	Comparison with the randomized method	18
8	Conclusion and Perspectives	19
	Appendices	21
	Appendix A Further results	21

Appendix B Reproduce the experiments	21
B.1 Packages and libraries	21
B.2 Configurations	21
Appendix C Run the project	23
C.1 Install the project	23
C.2 Set up conda environment:	23
C.2.1 Create the environment (might be slow):	23
C.2.2 Activate the environment:	23
C.2.3 Setup project:	23
C.3 Set up wandb for experiment tracking:	23
C.4 Train a model :	23
C.4.1 On a computer	23
C.4.2 On a cluster using SLURM	24
C.4.3 Configurations	24
C.5 Modify the code	25
C.6 Computing the metrics on the test data set	26
C.7 Compute the sign distance map	26
C.8 Perform the data augmentation	27
List of Figures	28
List of Tables	28

Abstract

Modeling the background error covariance matrix has a very positive impact on the quality of forecasts in numerical weather prediction. When using diffusion-based covariance operators, computing normalization coefficients to ensure that the diagonal elements of the modeled correlation matrix are all equal to one is a crucial component. In this paper, we study a deep learning approach with convolutional neural networks to estimate the normalization coefficients. Our approach is validated on the Nemovar data obtained by a brute-force exact method. We achieved good results with a maximum absolute relative error of 4.3% and a mean of 0.4% averaged over 190 samples. On a simplified Python framework, the CNN approach gives error metrics twice as low as the classically used randomized method.

1 Introduction

The aim of this study is to provide a proof of concept of the use of deep learning methods for the estimation of the normalization coefficients of diffusion operators in the assimilation of variational data. Current methods [1] for computing those coefficients are computationally costly, preventing the computation of the coefficient in each data assimilation cycle. The advantage of using neural networks (NN) to make inference is its low computational cost. However, the computation of the network weights during the training phase can be very expensive. So, neural networks are efficient methods for data assimilation allowing to take into account the temporal variations of the model coefficients with seasons and ocean currents. As shown in [2, 3, 4], we use the universality of a convolutional neural network (CNN) to approximate a continuous function with any accuracy, as long as the network is deep enough. Let f be the function we want to approximate.

$$\begin{aligned} f: \mathbb{R}^m \times \mathbb{R}^m \times \mathbb{R}^m &\rightarrow \mathbb{R}^m \\ (A_i, A_j, W) &\mapsto \Gamma \end{aligned}$$

with A_i (resp. A_j) the scaled diffusivity on axis y (resp. x), W the volume of the cells, Γ the normalization coefficients, m the dimension of the grid.

The input of the neural network is (A_i, A_j, W) which can be seen as a three-channel image.

By construction, the layers of a CNN are equivariant to translation [2]. This means that if the input of a function is translated, the output is translated in the same way, let f be a function and τ be a translation, we have $\tau \circ f(x) = f \circ \tau(x)$. For the function to approximate the equivariance to translation obtained by using the scaled diffusivity and the volume of the cells as described in [5].

2 Problem definition

The first goal is to determine whether a CNN can reach the performance of a U-Net without exceeding the number of trainable parameters. To do that, we have a pre-trained U-Net with one million parameters and a CNN on which to perform hyperparameter tuning. The hypothesis to test is that the better performance of the U-Net can be explained by the larger number of parameters alone.

Subsequently, the goal is to beat the randomized method [1] on the Nemovar data using a neural network.

To evaluate the performance, we compute the following metrics. Let us first define Γ the normalization coefficient, it is the target of the CNN, $\hat{\Gamma}$ the normalization coefficients estimated by the CNN, γ the vectorization of Γ noted $\gamma = \text{vec}(\Gamma)$ and \mathcal{O} the grid points in the oceans.

- For one sample $\Gamma \in \mathbb{R}^{h \times w}$,

1. The relative error $\varepsilon \in \mathbb{R}^{|\mathcal{O}|}$ is defined component-wise such as

$$\varepsilon_i = \frac{\hat{\gamma}_i^2 - \gamma_i^2}{\gamma_i^2}, \quad \forall i \in \mathcal{O}$$

2. RMSE(ε): it gives the average accuracy of the model, it is defined by $\text{RMSE}(\varepsilon) = \sqrt{\mathbb{E}[\varepsilon \circ \varepsilon]}$ with \circ the Hadamard product.
3. The quantile 99.99
4. The maximum relative error: $\varepsilon_{\max} = \max_{i \in \mathcal{O}} |\varepsilon_i|$
5. The mean absolute relative error: $\varepsilon_{\text{mean}} = \frac{1}{|\mathcal{O}|} \sum_{i \in \mathcal{O}} |\varepsilon_i|$

- For d samples $\Gamma^{(d)} : \mathbb{R}^{d \times h \times w}$

1. The relative error on a sample ℓ : $\varepsilon^{(\ell)}$ such as $\varepsilon_i^{(\ell)} = \frac{(\hat{\gamma}_i^{(\ell)})^2 - (\gamma_i^{(\ell)})^2}{(\gamma_i^{(\ell)})^2}, \forall i \in \mathcal{O}$,
with $\gamma^{(\ell)} = \text{vec}(\Gamma^{(\ell)})$
2. $\varepsilon_{\infty} : \frac{1}{d} \sum_{\ell=1}^d Q(|\varepsilon^{\ell}|, 0.9999)$
3. The maximum relative error: $\varepsilon_{\max} = \frac{1}{d} \sum_{\ell=1}^d \max_{i \in \text{Oceans}} |\varepsilon_i^{(\ell)}|$
4. The mean absolute relative error: $\varepsilon_{\text{mean}} = \frac{1}{d} \sum_{\ell=1}^d \frac{1}{|\mathcal{O}|} \sum_{i \in \mathcal{O}} |\varepsilon_i^{(\ell)}|$

We consider three datasets. The first one comes from a python toy model (Mercator projection) and the other two from Nemovar data (Partial_std, Full_std). Most of the work has been done on the Nemovar data.

In practice, each data set is divided into three sets for training, validation, and tests. For the Python data, we use 90 samples for training, 10 samples for validation, and none for the test. With the Nemovar datasets, there are 10 samples for training, 1 for validation, and 190 for the test. We used little data in training and validation for performance issues, as we wanted to reach 50,000 epochs to see if convergence was achieved. In addition, it is interesting to see if the method works using few data because if the research is to be extended to the three-dimensional case, there will be fewer data than in 2D.

For both python and Nemovar data, there is horizontal periodicity. However, for the Python data, there is no vertical periodicity because there are two continents on the north and south of the grid. On the other hand, there is vertical periodicity for the Nemovar data, and it depends on the characteristics of the grid; more details on the padding can be found in Folkes' report.

For the Nemovar data, the difference between the two data sets is based on how we standardized the data. For the first one, Partial_std, we only standardized the image using the coefficients in the ocean, while for the second one, Full_std, we standardized the whole image. The difference between the two datasets is enlightened in Figure 1, the color map is saturated but the important thing to notice is that for the partial standardization the value for the continental cells is 0 while for the oceanic cells the values are both positive and negative. However, for the full standardization, the value of the majority of oceanic pixels is positive. We use the first data set until Section 5.

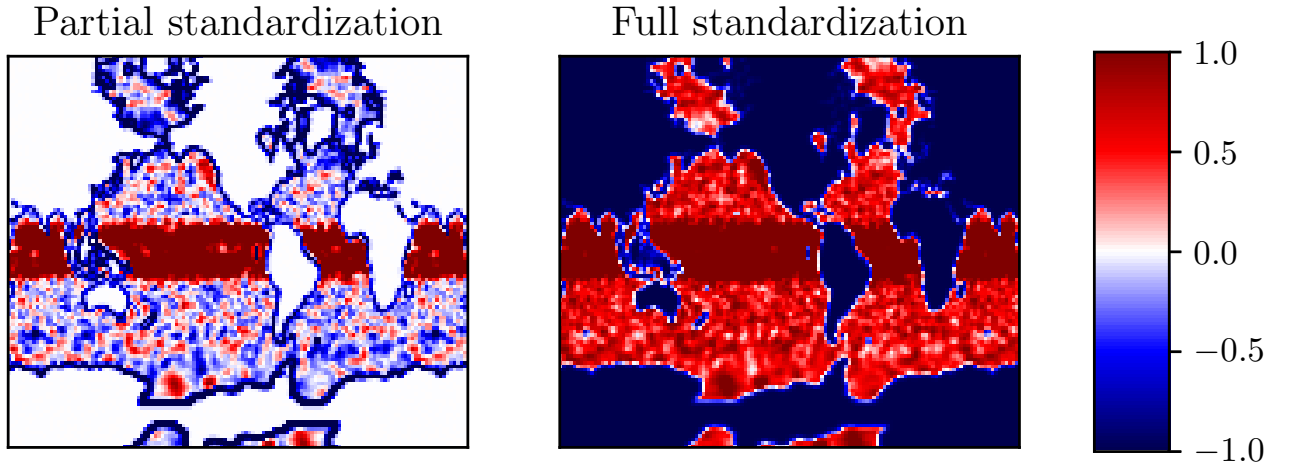


Figure 1: Difference between the two datasets for the first channel, A_i , of the input data

To evaluate the performance of data augmentation, we compare metrics and the localization of the relative error with and without data augmentation on the test dataset. The

test dataset consists of all samples that are not used in any training, in this case, there are 190 samples.

For the Nemovar datasets, each training is done on the same number of training step with number of training step = $\frac{\text{number of samples}}{\text{batch size}} \times \text{number of epochs}$, we fixed that number at 5×10^6 . This is useful as we will not have the same dataset length when using data augmentation.

3 Choosing the architecture

The goal is to determine if the CNN can beat the U-Net using the python data. We tried a CNN with several layers and 3×3 convolution, batch normalization, Rectified Linear Unit ($\text{ReLU} := \max(0, x)$) as an activation function for the layers and $\text{MSE}(\gamma)$ as a loss function. The different performances can be seen in Appendix A on Table 9, each metric is computed at the end of the 3000th epochs on the validation set. The best layer number is 10. This architecture has 3.36×10^5 trainable parameters.

Architecture	ε_{∞}	ε_{\max}	$\varepsilon_{\text{mean}}$
U-Net	7.8×10^{-2}	1.3×10^{-1}	4.3×10^{-3}
CNN	2.3×10^{-2}	4.8×10^{-2}	9.5×10^{-4}

Table 1: Performance of CNN and U-Net on python data

The architecture we retain is a CNN with 10 layers composed of:

- 1 input layer:
 - 3×3 convolution
 - 3 input channels, 64 output channels
 - Batch normalization
 - ReLU
- 8 hidden layers:
 - 3×3 convolution
 - 64 input channels, 64 output channels
 - Batch normalization
 - ReLU
- 1 output layer
 - 3×3 convolution

– 64 input channels, 1 output channels

Concerning the **padding** for the convolution, we used a padding *same*, meaning that the size of the image is conserved through the forward propagation, the *padding mode* is *replicate* meaning the input tensor is filled using the replication of the input limit. Later on the Nemovar data, we use padding *valid*, as we apply a 3×3 kernel for the convolution, the image size decreases by 1 for each layer.

Few words on batch normalization: as defined in PyTorch documentation¹ it consists in normalizing the data along the channels then multiplying the data by one vector and adding another one, both vectors being learnable parameters of the CNN. We compared the performance of this architecture with the dataset `Partial_std`, with and without batch normalization, and we obtain the best results with batch normalization. For the rest of this report, we will always apply batch normalization in the input and hidden layers. It is also important to consider how the data were initially standardized. For example, when using `Partial_std` as the normalization uses the continental cells that were not used for the standardization, it might induce a loss of information. This could explain the difference in performance between the two data sets in Section 5

4 Finding a good baseline

In this part, we try to find a good configuration of hyperparameters before modifying the data with data augmentation or changing the architecture with skip connections.

4.1 Distance map from coasts

As the errors were mostly located near the coast, we tried to add in the input a sign distance map from coasts to provide more information to the CNN. It is a sign distance map with the distance from the coast for the ocean cells and the opposite of the distance for the continental cells. We compare the performance using Manhattan distance and the Euclidean one with the initial performance. As shown in Table 2 the best result is achieved with the Manhattan distance. We mainly look at ε_{∞} and ε_{\max} as $\varepsilon_{\text{mean}}$ is already good, even if it is worse with the distance map. We trained a CNN on the data set `Partial_std` using 10 layers as described in Table 11, Experiment Distance map. For the rest of the report, we now consider Manhattan distance in input.

¹<https://pytorch.org/docs/stable/generated/torch.nn.BatchNorm2d.html>

	ε_{∞}	ε_{\max}	$\varepsilon_{\text{mean}}$
No map	7.29×10^{-2}	1.18×10^{-1}	4.33×10^{-3}
Manhattan	7.09×10^{-2}	1.12×10^{-1}	5.22×10^{-3}
Euclidean	7.37×10^{-2}	1.17×10^{-1}	4.82×10^{-3}

Table 2: Performance of CNN with different sign distance maps

4.2 Activation function

We compare different activation functions to the ReLU. The idea was to test smoother functions with the Exponential Linear Unit (ELU) and Softplus. We also try Parametric ReLU (PReLU) which is a Leaky ReLU (LReLU) but the α of the formulae given below is a trainable parameter; in practice, LReLU is often used.

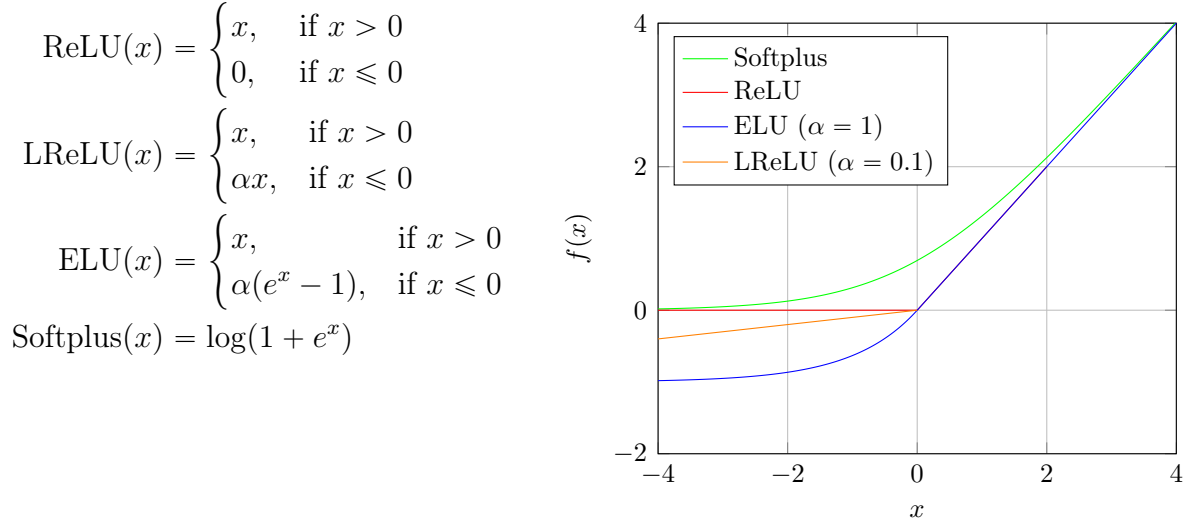


Figure 2: Different activation functions

In order to evaluate the performance, we trained the CNN on the same configuration as before but with the Manhattan distance as input, the complete configuration is described in Table 11, Experiment Activation function. ELU is the activation function that provides the best results, Softplus also performed well. But for the rest of the experiments, we will stick to the ELU.

Activation function	ε_{∞}	ε_{\max}	$\varepsilon_{\text{mean}}$
ReLU	7.09×10^{-2}	1.12×10^{-1}	5.22×10^{-3}
PLReLU	7.02×10^{-2}	1.10×10^{-1}	4.07×10^{-3}
ELU ($\alpha = 1$)	5.71×10^{-2}	8.86×10^{-2}	4.44×10^{-3}
Softmax	5.97×10^{-2}	9.38×10^{-2}	4.24×10^{-3}

Table 3: Performance of CNN with different activation function

The impact of the activation function is highly visible on the convergence of metrics in validation. In figure 3, we compare the convergence of loss, ε_{\max} , and $\varepsilon_{\text{mean}}$ in training and validation for the two CNNs with ELU and ReLU. With ReLU, a plateau is quickly reached for the validation. The training seems to have no impact on it, while with ELU, convergence is clearly better.




Figure 3: Difference in convergence using ELU and ReLU as activation functions

The CNN we now are trying to improve has an additional input with the sign distance map and has ELU as the activation function. This model has 2.99×10^5 trainable parameters.

5 Skip connection

One notorious issue of deep neural networks is the vanishing / exploding gradient, which is exacerbated by the number of layers causing degradation in accuracy. One way to overcome this problem is to use skip connections, this approach was first demonstrated in [6] with the ResNet network. Another way to address this issue is data normalization [7].

We tried to modify the architecture to see if it could improve the performance of our model. This section is divided into two parts, the first one concerns the data `Partial_std` and the second one the `Full_std`. In the first part, by mistake, the experiment was carried out using $\mathbb{E}[\varepsilon \circ \varepsilon]$ instead of $\text{MSE}(\gamma, \hat{\gamma})$. For time reasons, the impact of the loss function has only been evaluated on `Full_std` but not `Partial_std`.

5.1 First dataset

The first architecture we tried consisted in having blocks with two layers. Each block concatenates its input, i.e. the output of the previous block or layer, with the input of the previous block or layer. This architecture is described in Figure 4 with only two blocks. For 10 layers, with 4 blocks with 2 layers each, the model has 4.12×10^5 trainable parameter.




Figure 4: Architecture representation with 2 blocks and skip connections

Skip connections	ε_{∞}	ε_{\max}	$\varepsilon_{\text{mean}}$
Without	5.71×10^{-2}	8.86×10^{-2}	4.44×10^{-3}
With	4.39×10^{-2}	6.64×10^{-2}	3.89×10^{-3}

Table 4: Performance of CNN with and without skip connections

5.2 Change of dataset

In this part, the data used for the training and the test is now the dataset Full_std.

5.3 First architecture

We evaluate the performance of the previous model with different loss Experiment Loss and layers number Experiment Layers

Skip connections	loss	layers	ε_{∞}	ε_{\max}	$\varepsilon_{\text{mean}}$
Without	$\text{MSE}(\gamma, \hat{\gamma})$	10	2.99×10^{-2}	5.22×10^{-2}	4.07×10^{-3}
With	$\text{MSE}(\gamma, \hat{\gamma})$	10	2.83×10^{-2}	5.29×10^{-2}	4.32×10^{-3}
Without	$\mathbb{E}[\varepsilon \circ \varepsilon]$	10	2.25×10^{-2}	4.48×10^{-2}	3.72×10^{-3}
With	$\mathbb{E}[\varepsilon \circ \varepsilon]$	10	2.06×10^{-2}	4.24×10^{-2}	3.72×10^{-3}
With	$\mathbb{E}[\varepsilon \circ \varepsilon]$	16	2.06×10^{-2}	3.80×10^{-2}	3.55×10^{-3}

Table 5: Performance of CNN with and without skip connections with different loss functions and number of layers

The best results have been achieved with skip connections and $\mathbb{E}[\varepsilon \circ \varepsilon]$ as the loss function. In Figure 5, for the model with 10 layers, we can see that the worst performance, over 1%, is for pixels at a maximum distance of three from the coast. Furthermore, the further away the coast, the better the performance, which can be related to the distribution of ocean pixels by distance from the coast shown in Figure 6.

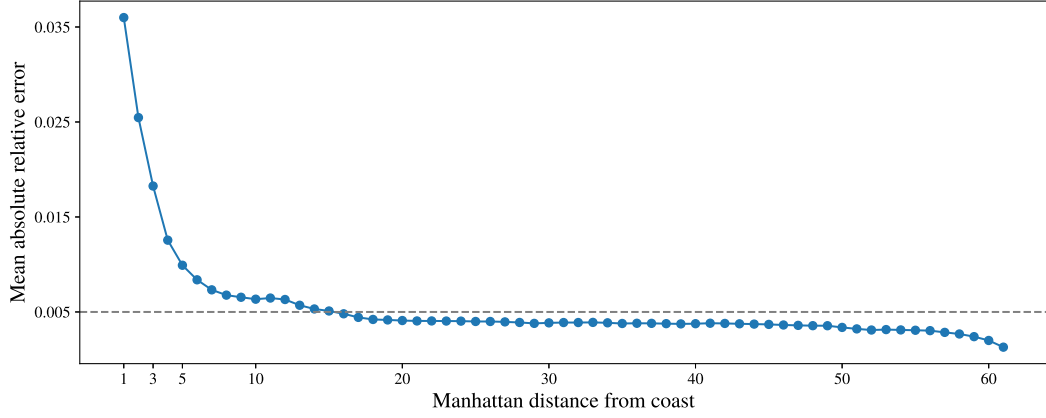


Figure 5: Mean absolute relative error with respect to the distance to the coast




Figure 6: Distribution of ocean pixels by distance from the coast

We see in Figure 7 the benefit of going to 16 layers is to reduce the maximum error for a distance of 1 to the coast. How with 16 layers the model has 7.45×10^5 trainable parameters against 4.12×10^5 and the training duration is 15h30 against 13 hours.

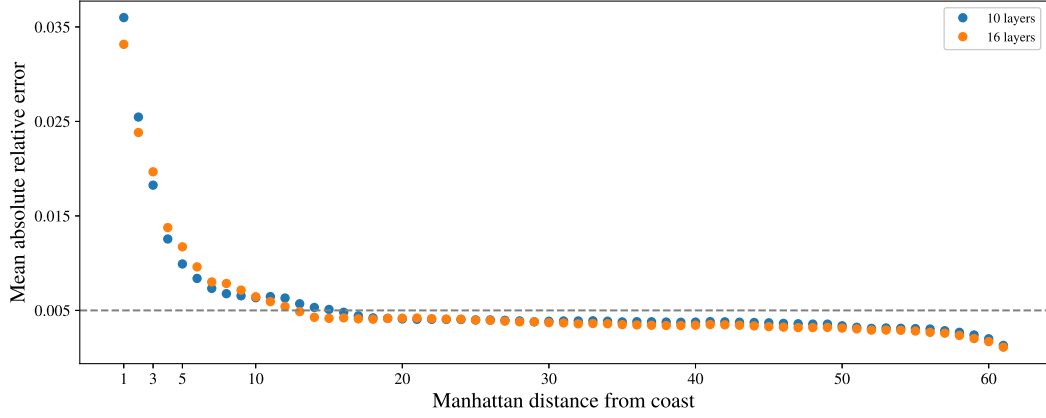


Figure 7: Comparison of the mean absolute relative error with respect to the distance to the coast between the best model with 10 and 16 layers

5.4 Second architecture

Afterward, we tried another type of block, instead of concatenating the data, each block adds up its input and output data as described in Figure 9. The CNN layout is shown in Figure 8. This architecture is inspired by the ResNet model [6].




Figure 8: Architecture representation with 2 blocks and skip connections

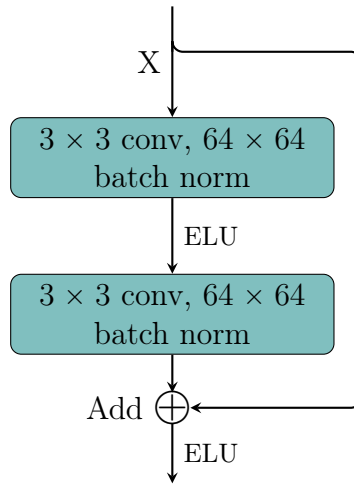


Figure 9: Block details

Skip connections	ε_{∞}	ε_{\max}	$\varepsilon_{\text{mean}}$
V1	2.06×10^{-2}	4.24×10^{-2}	3.72×10^{-3}
V2	1.87×10^{-2}	4.27×10^{-2}	3.80×10^{-3}

Table 6: Comparison of the performance using two different types of skip connection

For the same number of layers, the results are better for ε_{∞} with the second architecture but slightly worse for the others. However, this model has fewer trainable parameters than the first one, 2.99×10^5 against 4.12×10^5 . This allows a faster training speed: 9 hours instead of 13 hours to reach 5×10^4 epochs. We can see in Figure 10 that architecture 2 is better for pixels at distances 2 or 3 of the coast.




Figure 10: Comparison of the mean absolute relative error with respect to the distance to the coast between two types of skip connections




Figure 11: Map of $|\varepsilon|$ averaged on the test data set using the best model

6 Data augmentation

Data augmentation is commonly used for computer vision tasks with CNN, it aims to artificially increase the amount of data by transforming existing data. It helps to avoid overfitting by providing more data to the CNN [8], among the possible transformations we tried rotations and flips.

6.1 Method

As we have little data, we tried data augmentation with rotations and flips. Since the scale diffusivity α_i (resp. α_j) is defined on the north (resp. east) of the cell, the scale diffusivity needs to be changed according to the modification. For instance, with a vertical flip, if we look at a specific cell, the cell that was above is now below, meaning that the scale diffusivity α_i should be at the south as it is not defined we need to shift it by 1, as shown on Figure 12. For the 90° and 270° rotations, α_i and α_j are swapped as the axes are rotated. As the models' inputs are standardized, α_i and α_j must be properly restandardized.




Figure 12: Representation of the vertical flip




Figure 13: Representation of the 90° rotation

6.2 Results

We performed the training on the model with skip connections, $\mathbb{E}[\varepsilon \circ \varepsilon]$ as loss function, and 10 layers as described in Experiment Augmentation. Augmentation is made on the dataset Full_std.

Data set	ε_{∞}	ε_{\max}	$\varepsilon_{\text{mean}}$
Original (10/1)	2.06×10^{-2}	4.24×10^{-2}	3.72×10^{-3}
Augmented with horizontal flips (20/1)	2.27×10^{-2}	4.78×10^{-2}	3.74×10^{-4}

Table 7: Impact of data augmentation.

The other transformations led to poor results because the model converges quickly in validation, so training has no impact on validation.

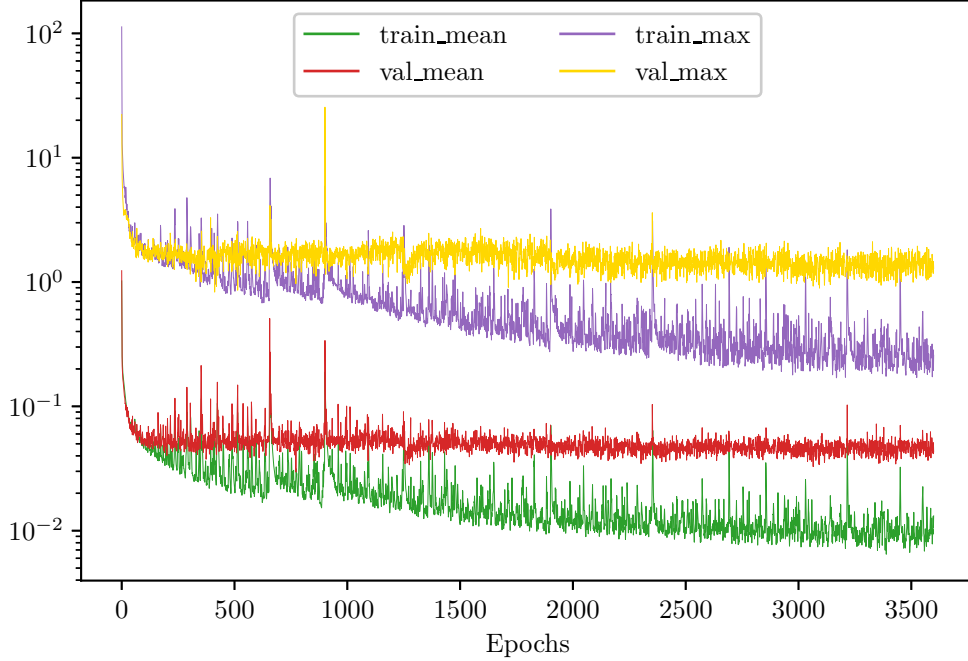


Figure 14: Difference in convergence in training and validation with augmented data with 90° rotation

7 Comparison with the randomized method

In order to compare the performance of the CNN with the randomized method with 10,000 samples on the Python data. We apply the best model of the CNN, which we found by training on Nemovar data, on the Python data. The problem is slightly different between the python and Nemovar data, but this gives a good idea of the potential of the approach. The randomized method is described in [1]. As we can see in Table 8 the metrics are better by a factor of 2 for the CNN. But it must be noticed for the CNN ε_{\max} is higher as the problem is more difficult, but ε_{∞} is similar.

Method	ε_{∞}	ε_{\max}	$\varepsilon_{\text{mean}}$
Randomized (10^4)	5.4×10^{-2}	5.6×10^{-2}	1.1×10^{-2}
CNN	1.9×10^{-2}	2.8×10^{-2}	6.5×10^{-3}

Table 8: Comparison of randomization and deep learning

8 Conclusion and Perspectives

We achieved really good results with the CNN using skip connections and Exponential Linear Unit, the error is still high near the coasts but further tuning can be done to improve the performance. But the goal of showing a proof of concept is reached and, in the context of the experiment, CNN performed better than randomization.

To go further with the CNN, we thought of doing the training on sections of the image instead of the whole image, the interest that allows trying different strategies in the selection of the sections. It could be useful for importance sampling, but also for the computation of the training data, as only values for some pixels are needed. The major drawback of this approach is the computation cost, as there might be thousands of sections for a single image, making data loading a bottleneck.

Finally, the main issue is to move to the three-dimensional case, the first clues would be to either use 3d convolution kernels or to apply the CNN to one layer and provide information on the other layers. However, the critical issue will be the computation cost of the true coefficients and the training.

References

- [1] Anthony T. Weaver, Marcin Chrust, Benjamin Ménétrier, and Andrea Piacentini. An evaluation of methods for normalizing diffusion-based covariance operators in variational data assimilation. *Quarterly Journal of the Royal Meteorological Society*, 147 (734):289–320, 2021. URL <https://onlinelibrary.wiley.com/doi/abs/10.1002/qj.3918>.
- [2] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. URL <http://www.deeplearningbook.org>.
- [3] Ding-Xuan Zhou. Universality of Deep Convolutional Neural Networks, July 2018. URL <http://arxiv.org/abs/1805.10769>.
- [4] Philipp Petersen and Felix Voigtlaender. Equivalence of approximation by convolutional neural networks and fully-connected networks. January 2021. URL <http://arxiv.org/abs/1809.00973>.
- [5] Olivier Goux. Correlation models based on diffusion. Internal note, 2022.
- [6] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep Residual Learning for Image Recognition, December 2015. URL <http://arxiv.org/abs/1512.03385>.
- [7] Bouman Kak. Using skip connections to mitigate the problem of vanishing gradients, and using batch, instance, and layer normalizations for improved sgd in deep networks. Purdue University, April 2022. URL <https://engineering.purdue.edu/DeepLearn/pdf-kak/SkipConsAndBN.pdf>.

- [8] Connor Shorten and Taghi M. Khoshgoftaar. A survey on Image Data Augmentation for Deep Learning. *Journal of Big Data*, 6(1):60, July 2019. URL <https://doi.org/10.1186/s40537-019-0197-0>.

Appendices

Appendix A Further results

Layers	Parameters	ε_{∞}	ε_{\max}	$\varepsilon_{\text{mean}}$
5	1.5×10^5	3.1×10^{-2}	7.1×10^{-2}	1.2×10^{-3}
10	3.4×10^5	2.3×10^{-2}	4.8×10^{-2}	9.5×10^{-4}
15	5.2×10^5	4.4×10^{-2}	9.7×10^{-2}	1.8×10^{-3}
20	7.0×10^5	4.2×10^{-2}	9.1×10^{-2}	1.7×10^{-3}

Table 9: Performance of CNN for several layers number on python data

Appendix B Reproduce the experiments

B.1 Packages and libraries

python: 3.9.12
conda: 4.13.0
conda-build: 3.21.9

Package	Version	Channel
pytorch	1.11.0	pytorch
pytorch-lightning	1.6.4	conda-forge
torchvision	0.12.0	pytorch
wandb	0.12.17	conda-forge
numpy	1.22.3	

Table 10: Package versions

B.2 Configurations

Experiment	Distance map	Loss	layers	Channels	Activation functions	Data	Skip connection	Block/layer per block
Distance map	Var	$\text{MSE}(\gamma, \hat{\gamma})$	10	64	ReLU	Partial_std	None	
Activation function	Manhattan	$\text{MSE}(\gamma, \hat{\gamma})$	10	64	Var	Partial_std	None	
Skip connections	Manhattan	$\mathbb{E}[\varepsilon \circ \varepsilon]$	10	64	ELU	Partial_std	Var	4/2
Loss	Manhattan	Var	10	64	ELU	Full_std	With	4/2
Layers	Manhattan	$\mathbb{E}[\varepsilon \circ \varepsilon]$	Var	64	ELU	Full_std	With	Var /2
Augmentation	Manhattan	$\mathbb{E}[\varepsilon \circ \varepsilon]$	10	64	ELU	Var	With	4/2

Table 11: Configurations of the different experiments

Appendix C Run the project

C.1 Install the project

```
git clone https://github.com/FolkeKS/DL-normalization.git
```

C.2 Set up conda environment:

C.2.1 Create the environment (might be slow):

```
conda env create --file environment.yml
```

- If the installation is stuck on Solving environment: | try:

```
conda config --set channel_priority strict
```

- Revert with:

```
conda config - --set channel_priority true
```

C.2.2 Activate the environment:

```
bash conda activate DL-normalization
```

C.2.3 Setup project:

```
bash pip install -e .
```

C.3 Set up wandb for experiment tracking:

- Sign up at <https://wandb.ai/site> and log in
- Find your API-key at <https://wandb.ai/authorize>
- With your conda environment activated, run the following command and provide API-key

```
wandb login
```

C.4 Train a model :

C.4.1 On a computer

```
python scripts/trainer.py fit --config configs/demo.yaml
```

C.4.2 On a cluster using SLURM

```
sbatch scripts/train.bash
```

C.4.3 Configurations

The configuration can be modified in the yaml files in [configs/](#). The model part, shown below, describes the parameter to instantiate the class CNN in [src/cnn.py](#).

```
n_blocks: 4
n_blocks_filters: 64
layers_per_block: 2
kernel_size: 3
n_channels: 4
n_classes: 1
q: 0.9999
standarize_outputs: true
predict_squared: false
predict_inverse: false
loss_fn: masked_mse
padding_type: "valid"
optimizer: Adam
data_dir: data/processed/sections/
```

This is the equivalent to

```
class CNN(pl.LightningModule):
    def __init__(self,
        n_blocks: int = 4,
        n_blocks_filters: int = 64,
        layers_per_block: int = 2,
        kernel_size: int = 3,
        n_channels: int = 4,
        n_classes: int = 1,
        q: float = 0.9999,
        standarize_outputs: bool = False,
        predict_squared: bool = False,
        predict_inverse: bool = False,
        loss_fn: str = "masked_mse",
        padding_type: str = "valid",
        optimizer: str = "Adam",
```



```
data_dir: str = "data/processed/newdata/",
**kwargs):
```

The data part describes the parameter to instantiate the class `DirLightDataset` in [src/data/dataset.py](#)

```
class DirLightDataset(pl.LightningDataModule):
    def __init__(self,
        batch_size: int = 1,
        data_dir: str = "data/processed/newdata/",
        num_workers: int = 0,
        gpus: int = 0):
```

- In order to make deterministic runs we set `seed_everything: 42` (line 1) and `deterministic: true` (line 78).
- To continue a run after it has been stopped we set `ckpt_path: "results/wandb/cnn/36c5vu01/"` (line 115). With "cnn/" the wandb project in which the run was created, and "36c5vu01" the 'id' of the run. We also set `version: 36c5vu01` (line 11) to make wandb continue the training in the run instead of creating a new run. The `global_step` variable will be reset, so the charts must be visualized with `epoch` or `step` as x-axis.
- It is also possible to choose how the best checkpoint is selected (lines 20-35), we choose the checkpoint that minimizes the loss in validation, but other choices are possible as choosing the one minimizing the quantile. However, we have not found how to use two strategies for the same run.

C.5 Modify the code

The model is loaded in `scripts/trainer.py`. Modifications are to be done in the config file but there are also modifications to make in the module loading the data in [src/data/dataset.py](#):

- adjust the padding, by default the padding is 31 for the latitude and 28 for the longitude. The images are cropped on the fly, the dimension of the image taken by the CNN is $4 \times 292 + 2\ell \times 360 + 2\ell$ with ℓ the number of layers
- adding the sign distance map

For computational time, these modifications should be made to the data directly before starting the training.

```
class DirDataset(Dataset):
    def __getitem__(self, i):
        idx = self.ids[i]
        X_files = glob.glob(os.path.join(self.X_dir, idx+'.*'))
```

```

        Y_files = glob.glob(os.path.join(self.Y_dir, idx+'*_norm_coeffs.*'))
    #Load the input / true data
    X = torch.from_numpy(np.load(X_files[0])['arr_0']).float()
    Y = torch.from_numpy(np.load(Y_files[0])['arr_0']).float()
    #Load the distance map
    distance_map = np.load("data/python_sign_dist_map_std.npz")['arr_0']
    distance_map = torch.from_numpy(distance_map).float()
    #Crop the input data
    distance_map = transforms.CenterCrop([200, 360+2*10])(distance_map)
    X = transforms.CenterCrop([200, 360+2*10])(X)
    #Add the distance map
    X = torch.cat((X,torch.unsqueeze(distance_map, 0)),0)
    return X, \
        Y

```

C.6 Computing the metrics on the test data set

As explained in the report, we used one data set for the python data two data sets **Partial std** = **finaldata** and **Full std** = **newdata** for the Nemovar data. However, for the **Partial std** data set the training data were standardized using 10 samples while the test data were standardized using 180 samples. So, to compute the metrics the test data had to be destandardized and then restandardized. For **Full std** all the data are standarized using the 10 training samples. We compute the mean and standard deviation of the mean/max/quantile 99,99% of the absolute relative error over the test dataset. We save the tensor of the relative error for each sample and save an image of the mean of the relative error. The code and results are in the repository [results/test_metrics/](#) each repository corresponds to a data set. To run the computation of the metrics for the root directory of the project the python script. The configuration of the runs is added to a list **model_params**, each element of this list is a list with the index corresponding to:

0. Boolean: compute or not the metrics
1. Boolean: if the run uses the sign distance map
2. Numpy array: the sign distance map
3. Int: number of layers
4. Class: the model loaded with the correct src

C.7 Compute the sign distance map

The sign distance map is computed in the notebook [notebooks/Distance_map.ipynb](#). We use the method `scipy.ndimage.distance_transform_bf` to compute the distance ([doc link](#)).

C.8 Perform the data augmentation

The script to use is [src/data/augmentation.py](#) and the combinations of transformations are tested in the notebook [notebooks/test_augmentation.ipynb](#). As the sign distance map is added to the input data during the augmentation, it is important to not add it another time as described in Section C.5.

List of Figures

1	Difference between the two datasets for the first channel, A_i , of the input data	5
2	Different activation functions	8
3	Difference in convergence using ELU and ReLU as activation functions	9
4	Architecture representation with 2 blocks and skip connections	10
5	Mean absolute relative error with respect to the distance to the coast	12
6	Distribution of ocean pixels by distance from the coast	12
7	Comparison of the mean absolute relative error with respect to the distance to the coast between the best model with 10 and 16 layers	13
8	Architecture representation with 2 blocks and skip connections	14
9	Block details	14
10	Comparison of the mean absolute relative error with respect to the distance to the coast between two types of skip connections	15
11	Map of $ \varepsilon $ averaged on the test data set using the best model	16
12	Representation of the vertical flip	17
13	Representation of the 90° rotation	17
14	Difference in convergence in training and validation with augmented data with 90° rotation	18

List of Tables

1	Performance of CNN and U-Net on python data	6
2	Performance of CNN with different sign distance maps	8
3	Performance of CNN with different activation function	9
4	Performance of CNN with and without skip connections	11
5	Performance of CNN with and without skip connections with different loss functions and number of layers	11
6	Comparison of the performance using two different types of skip connection	15
7	Impact of data augmentation.	17
8	Comparison of randomization and deep learning	18
9	Performance of CNN for several layers number on python data	21
10	Package versions	21
11	Configurations of the different experiments	22